


# 200行代码讲透Rust Futures的问题

---

 [stevenbai.top/rust/futures\\_explained\\_in\\_200\\_lines\\_of\\_rust2](https://stevenbai.top/rust/futures_explained_in_200_lines_of_rust2)

rust

2020/4/13

---

## 引言

---

在翻译完本文以后很快就有人反馈作者实现的runtime有问题,然后作者进行了修复,没想到发下了更大的问题. 当然最终解决了. 我们来看看过程,同时加深对Future的理解.

作者关于此问题的讨论过程见github issue.

本文是在有问题的版本基础之上撰写,最终修订后,正常的版本在这里.

## 问题1: Future返回Pending一定要有唤醒机制

---

将main函数改写如下:

```

fn main() {
    let start = Instant::now();
    let reactor = Reactor::new();
    let reactor = Arc::new(Mutex::new(reactor));
    let future1 = Task::new(reactor.clone(), 1, 1);
    let future2 = Task::new(reactor.clone(), 2, 2);

    let start1 = start.clone();
    let fut1 = async move {
        let val = future1.await;
        let dur = (Instant::now() - start1).as_secs_f32();
        println!("Future got {} at time: {:.2}.", val, dur);
    };

    let start2 = start.clone();
    let fut2 = async move {
        let val = future2.await;
        let dur = (Instant::now() - start2).as_secs_f32();
        println!("Future got {} at time: {:.2}.", val, dur);
    };

    let mainfut = async {
        fut1.await;
        fut2.await;
    };

    println!("Running on first executor");
    let mainfut = match block_on(futures::future::select(
        Box::pin(mainfut),
        futures::future::ready(()),
    )) {
        futures::future::Either::Left(_) => unreachable!(),
        futures::future::Either::Right(((), mainfut)) => mainfut,
    };

    println!("Running on second executor");
    std::thread::spawn(move || {
        block_on(mainfut);
    })
    .join();

    reactor.lock().map(|mut r| r.close()).unwrap();
}

```

## Rust

然后尝试运行,你会发现第40行的join无法返回,为什么呢?

我们来分析一下task1的运行过程: 1. select的那次poll,这次因为没有注册,所以通过channel 发送 **Event::Timeout** 给Reactor,Reactor则会启动一个线程T1运行该任务. 2. select这个future会提前结束,并没有等待刚刚创建的线程T1结束 3. 再次创建一个新的线程T2来运行task1,这时候task1的状态是已经注册,所以它就直接返回pending了.

然后,就再也不能唤醒了,为什么呢,因为task1返回了Pending,但是没有任何唤醒当前线程的机制T2,因为T1结束后唤醒的是T1,而不是T2.

## 修复Future唤醒问题

---

知道了问题所在,相应的也就容易得多. 这也是赖智超的博客中说明的问题.

这里转述如下: > 在Rust的异步模型中, 有一个极其关键的细节容易被忽略: 在多次调用Future::poll方法时, Executor传递进去的Waker有可能是不一样的, 因此每次返回Pending之前都需要把之前保存的Waker更新. 比如考虑一个带工作窃取的Executor, Future刚开始在T1线程执行, 注册了Waker, 由于T1负载过高, Future被T2线程窃取执行, 这时如果不更新Waker, 当Future真正ready的时候, 调用老的Waker, 通知的却是T1线程, 导致真正拥有Future的T2丢失了通知, Future永远无法结束.

那么解决方法的关键就是每次返回Pending的时候,保证Waker唤醒的线程就是block\_on所在的那个线程即可.

这主要涉及到三个地方,下面分别描述.

### Reactor的改动1

---

Reactor要保存每个task id对应的Waker,这个和初次运行Event::Timeout传递进来的Waker可能不是同一个. 所以Reactor的声明如下:

```
struct Reactor {
    dispatcher: Sender<Event>,
    handle: Option<JoinHandle<()>>,
    readylist: Arc<Mutex<Vec<usize>>>,
    wakers: Arc<Mutex<HashMap<usize, Waker>>>,
}
#[derive(Debug)]
enum Event {
    Close,
    Timeout(u64, usize),
}
```

Rust

注意到区别: 1. 多了一个wakers字段,其中的key是任务id,value则是关联的waker 2. Event的Timeout不再传递Waker,只是传递任务id

### Reactor的改动2

---

说白了就是需要将任务的id和Waker关联起来,这样真正运行任务的线程结束后,可以调用相应的wake函数.

```

fn insert_or_swap(&mut self, id: usize, waker: Waker) {
    self.wakers
        .lock()
        .map(|mut wakers| {
            let _ = wakers.insert(id, waker);
        })
        .unwrap();
}

fn register(&mut self, duration: u64, waker: Waker, data: usize) {
    self.insert_or_swap(data, waker);
    self.dispatcher
        .send(Event::Timeout(duration, data))
        .unwrap();
}

```

Rust

## Task的poll实现

---

```

fn poll(mut self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output> {
    println!(
        "poll task {}, is_registered={},thread={:?}",
        self.id,
        self.is_registered,
        thread::current().id(),
    );
    let mut r = self.reactor.lock().unwrap();
    if r.is_ready(self.id) {
        Poll::Ready(self.id)
    } else if self.is_registered {
        r.insert_or_swap(self.id, cx.waker().clone());
        Poll::Pending
    } else {
        r.register(self.data, cx.waker().clone(), self.id);
        drop(r);
        self.is_registered = true;
        Poll::Pending
    }
}

```

Rust

最关键的地方就是这里:

```

if self.is_registered {
    r.insert_or_swap(self.id, cx.waker().clone());
    Poll::Pending
}

```

Rust

这样保证每次返回Pending,都是和当前的waker关联起来,因为这有可能发生变化.

## 休眠唤醒问题

---

如果我们仔细深入思考,这里其实还存在一个问题,就是当poll正在运行的时候被唤醒的问题.也就是

```
r.register(self.data, cx.waker().clone(), self.id);
    drop(r);
    //--->这里有竞争冲突
    self.is_registered = true;
```

### Rust

当然这里不是传统意义上的内存共享问题,更确切来说应该是死锁问题,这里可能reactor中很快就已经结束了任务的执行,然后尝试去唤醒executor线程,但是这时候实际上executor线程正在运行着.然后poll返回pending,executor线程进入休眠.

如果你用代码实验,就会发现不会发生死锁.这是因为thread::park文当中写明了,先unpark后park,这时候park会立即返回.但是如果你用其他导致线程休眠唤醒的方式就会出问题.这里说的是这种情况应该避免.

## 潜在死锁问题的解决办法

---

知道了问题所在,解决起来就会容易.实际上并发问题,难的是找到问题,理解问题.大多数时候解决起来反而容易.

问题的本质是:我们在executor在执行poll的过程,可能会被wake,但是这时候执行线程还没有休眠.

解决办法则是: 1. reactor线程唤醒的时候要持有reactor锁.

```

fn new() -> Arc<Mutex<Box<Self>>> {
    let (tx, rx) = channel::<Event>();
    let reactor = Arc::new(Mutex::new(Box::new(Reactor {
        dispatcher: tx,
        handle: None,
        tasks: HashMap::new(),
    })));

    let reactor_clone = Arc::downgrade(&reactor);
    let handle = thread::spawn(move || {
        let mut handles = vec![];
        for event in rx {
            let reactor = reactor_clone.clone();
            match event {
                Event::Close => break,
                Event::Timeout(duration, id) => {
                    let event_handle = thread::spawn(move || {
                        thread::sleep(Duration::from_millis(duration));
                        let reactor = reactor.upgrade().unwrap();
                        reactor.lock().map(|mut r| r.wake(id)).unwrap();
                        println!("wake {}", id);
                    });
                    handles.push(event_handle);
                }
            }
        }
        handles
            .into_iter()
            .for_each(|handle| handle.join().unwrap());
    });
    reactor.lock().map(|mut r| r.handle = Some(handle)).unwrap();
    reactor
}

```

Rust

## 1. poll的过程中也要持有reactor的锁.

```
impl Future for Task {
    type Output = usize;
    fn poll(self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output> {
        let mut r = self.reactor.lock().unwrap();
        if r.is_ready(self.id) {
            println!("POLL: TASK {} IS READY", self.id);
            *r.tasks.get_mut(&self.id).unwrap() = TaskState::Finished;
            Poll::Ready(self.id)
        } else if r.tasks.contains_key(&self.id) {
            println!("POLL: REPLACED WAKER FOR TASK: {}", self.id);
            r.tasks
                .insert(self.id, TaskState::NotReady(cx.waker().clone()));
            Poll::Pending
        } else {
            println!(
                "POLL: REGISTERED TASK: {}, WAKER: {:?}" ,
                self.id,
                cx.waker()
            );
            r.register(self.data, cx.waker().clone(), self.id);
            Poll::Pending
        }
    }
}
```

Rust

但是我不认为这个完整解决了问题,还是有可能发生唤醒发生在休眠之前的问题.

更好的方式是负责休眠/唤醒顺序的是executor,而不是具体的Future. 比如  
<https://github.com/stjepang/byo-executor/blob/master/examples/v1.rs>.

## 相关源码完整版

---

## 问题1的完整代码

---

```

use futures::future::select;
// from https://cfsamson.github.io/books-futures-explained/8_finished_example.html

use std::{
    future::Future,
    pin::Pin,
    sync::{
        mpsc::{channel, Sender},
        Arc, Mutex,
    },
    task::{Context, Poll, RawWaker, RawWakerVTable, Waker},
    thread::{self, JoinHandle},
    time::{Duration, Instant},
};

fn main() {
    let start = Instant::now();
    let reactor = Reactor::new();
    let reactor = Arc::new(Mutex::new(reactor));
    let future1 = Task::new(reactor.clone(), 1, 1);
    let future2 = Task::new(reactor.clone(), 2, 2);

    let start1 = start.clone();
    let fut1 = async move {
        let val = future1.await;
        let dur = (Instant::now() - start1).as_secs_f32();
        println!("Future got {} at time: {:.2}.", val, dur);
    };

    let start2 = start.clone();
    let fut2 = async move {
        let val = future2.await;
        let dur = (Instant::now() - start2).as_secs_f32();
        println!("Future got {} at time: {:.2}.", val, dur);
    };

    let mainfut = async {
        fut1.await;
        fut2.await;
    };

    println!("Running on first executor");
    let mainfut = match block_on(futures::future::select(
        Box::pin(mainfut),
        futures::future::ready(()),
    )) {
        futures::future::Either::Left(_) => unreachable!(),
        futures::future::Either::Right(((), mainfut)) => mainfut,
    };

    println!("Running on second executor");
    std::thread::spawn(move || {
        block_on(mainfut);
    })
    .join();
}

```



```

        reactor.lock().map(|mut r| r.close()).unwrap();
    }

// ===== EXECUTOR =====
fn block_on<F: Future>(mut future: F) -> F::Output {
    let mywaker = Arc::new(MyWaker {
        thread: thread::current(),
    });
    let waker = waker_into_waker(Arc::into_raw(mywaker));
    let mut cx = Context::from_waker(&waker);

    // SAFETY: we shadow `future` so it can't be accessed again.
    let mut future = unsafe { Pin::new_unchecked(&mut future) };
    let val = loop {
        match Future::poll(future.as_mut(), &mut cx) {
            Poll::Ready(val) => break val,
            Poll::Pending => thread::park(),
        };
    };
    val
}

// ===== FUTURE IMPLEMENTATION =====
#[derive(Clone)]
struct MyWaker {
    thread: thread::Thread,
}

#[derive(Clone)]
pub struct Task {
    id: usize,
    reactor: Arc<Mutex<Reactor>>,
    data: u64,
    is_registered: bool,
}

fn mywaker_wake(s: &MyWaker) {
    let waker_ptr: *const MyWaker = s;
    let waker_arc = unsafe { Arc::from_raw(waker_ptr) };
    println!("wake thread {:?}", waker_arc.thread.id());
    waker_arc.thread.unpark();
}

fn mywaker_clone(s: &MyWaker) -> RawWaker {
    let arc = unsafe { Arc::from_raw(s) };
    std::mem::forget(arc.clone()); // increase ref count
    RawWaker::new(Arc::into_raw(arc) as *const (), &VTABLE)
}

const VTABLE: RawWakerVTable = unsafe {
    RawWakerVTable::new(
        |s| mywaker_clone(&(s as *const MyWaker)), // clone
        |s| mywaker_wake(&(s as *const MyWaker)), // wake
        |s| mywaker_wake(*(&(s as *const &MyWaker))), // wake by ref
    )
}

```

```

        |s| drop(Arc::from_raw(s as *const MyWaker)), // decrease refcount
    )
};

fn waker_into_waker(s: *const MyWaker) -> Waker {
    let raw_waker = RawWaker::new(s as *const (), &VTABLE);
    unsafe { Waker::from_raw(raw_waker) }
}

impl Task {
    fn new(reactor: Arc<Mutex<Reactor>>, data: u64, id: usize) -> Self {
        Task {
            id,
            reactor,
            data,
            is_registered: false,
        }
    }
}

impl Future for Task {
    type Output = usize;
    fn poll(mut self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output>
    {
        println!(
            "poll task {}, is_registered={},thread={:?}",
            self.id,
            self.is_registered,
            thread::current().id(),
        );
        let mut r = self.reactor.lock().unwrap();
        if r.is_ready(self.id) {
            Poll::Ready(self.id)
        } else if self.is_registered {
            Poll::Pending
        } else {
            r.register(self.data, cx.waker().clone(), self.id);
            drop(r);
            self.is_registered = true;
            Poll::Pending
        }
    }
}

// ===== REACTOR =====
struct Reactor {
    dispatcher: Sender<Event>,
    handle: Option<JoinHandle<()>>,
    readylist: Arc<Mutex<Vec<usize>>>,
}

#[derive(Debug)]
enum Event {
    Close,
    Timeout(Waker, u64, usize),
}

```

```

impl Reactor {
    fn new() -> Self {
        let (tx, rx) = channel::<Event>();
        let readylist = Arc::new(Mutex::new(vec![]));
        let rl_clone = readylist.clone();
        let mut handles = vec![];
        let handle = thread::spawn(move || {
            // This simulates some I/O resource
            for event in rx {
                println!("REACTOR: {:?}", event);
                let rl_clone = rl_clone.clone();
                match event {
                    Event::Close => break,
                    Event::Timeout(waker, duration, id) => {
                        let event_handle = thread::spawn(move || {
                            thread::sleep(Duration::from_secs(duration));
                            rl_clone.lock().map(|mut rl| rl.push(id)).unwrap();
                            waker.wake();
                        });
                        handles.push(event_handle);
                    }
                }
            }

            for handle in handles {
                handle.join().unwrap();
            }
        });

        Reactor {
            readylist,
            dispatcher: tx,
            handle: Some(handle),
        }
    }

    fn register(&mut self, duration: u64, waker: Waker, data: usize) {
        self.dispatcher
            .send(Event::Timeout(waker, duration, data))
            .unwrap();
    }

    fn close(&mut self) {
        self.dispatcher.send(Event::Close).unwrap();
    }

    fn is_ready(&self, id_to_check: usize) -> bool {
        self.readylist
            .lock()
            .map(|rl| rl.iter().any(|id| *id == id_to_check))
            .unwrap()
    }
}

```

```
impl Drop for Reactor {  
    fn drop(&mut self) {  
        self.handle.take().map(|h| h.join().unwrap()).unwrap();  
    }  
}
```

Rust

## 修复Future唤醒问题的完整源码

---

```

use futures::future::select;
// from https://cfsamson.github.io/books-futures-explained/8_finished_example.html

use std::collections::HashMap;
use std::{
    future::Future,
    pin::Pin,
    sync::{
        mpsc::{channel, Sender},
        Arc, Mutex,
    },
    task::{Context, Poll, RawWaker, RawWakerVTable, Waker},
    thread::{self, JoinHandle},
    time::{Duration, Instant},
};

fn main() {
    let start = Instant::now();
    let reactor = Reactor::new();
    let reactor = Arc::new(Mutex::new(reactor));
    let future1 = Task::new(reactor.clone(), 1, 1);
    let future2 = Task::new(reactor.clone(), 2, 2);

    let start1 = start.clone();
    let fut1 = async move {
        let val = future1.await;
        let dur = (Instant::now() - start1).as_secs_f32();
        println!("Future got {} at time: {:.2}.", val, dur);
    };

    let start2 = start.clone();
    let fut2 = async move {
        let val = future2.await;
        let dur = (Instant::now() - start2).as_secs_f32();
        println!("Future got {} at time: {:.2}.", val, dur);
    };

    let mainfut = async {
        fut1.await;
        fut2.await;
    };

    println!("Running on first executor");
    let mainfut = match block_on(futures::future::select(
        Box::pin(mainfut),
        futures::future::ready(()),
    )) {
        futures::future::Either::Left(_) => unreachable!(),
        futures::future::Either::Right(((), mainfut)) => mainfut,
    };

    println!("Running on second executor");
    std::thread::spawn(move || {
        block_on(mainfut);
    })
}

```

```

        .join();

    reactor.lock().map(|mut r| r.close()).unwrap();
}

// ===== EXECUTOR =====
fn block_on<F: Future>(mut future: F) -> F::Output {
    let mywaker = Arc::new(MyWaker {
        thread: thread::current(),
    });
    let waker = waker_into_waker(Arc::into_raw(mywaker));
    let mut cx = Context::from_waker(&waker);

    // SAFETY: we shadow `future` so it can't be accessed again.
    let mut future = unsafe { Pin::new_unchecked(&mut future) };
    let val = loop {
        match Future::poll(future.as_mut(), &mut cx) {
            Poll::Ready(val) => break val,
            Poll::Pending => thread::park(),
        };
    };
    val
}

// ===== FUTURE IMPLEMENTATION =====
#[derive(Clone)]
struct MyWaker {
    thread: thread::Thread,
}

#[derive(Clone)]
pub struct Task {
    id: usize,
    reactor: Arc<Mutex<Reactor>>,
    data: u64,
    is_registered: bool,
}

fn mywaker_wake(s: &MyWaker) {
    let waker_ptr: *const MyWaker = s;
    let waker_arc = unsafe { Arc::from_raw(waker_ptr) };
    println!("wake thread {:?}", waker_arc.thread.id());
    waker_arc.thread.unpark();
}

fn mywaker_clone(s: &MyWaker) -> RawWaker {
    let arc = unsafe { Arc::from_raw(s) };
    std::mem::forget(arc.clone()); // increase ref count
    RawWaker::new(Arc::into_raw(arc) as *const (), &VTABLE)
}

const VTABLE: RawWakerVTable = unsafe {
    RawWakerVTable::new(
        |s| mywaker_clone(&(s as *const MyWaker)), // clone
        |s| mywaker_wake(&(s as *const MyWaker)),  // wake
    )
}

```

```

        |s| mywaker_wake(*(s as *const &MyWaker)),    // wake by ref
        |s| drop(Arc::from_raw(s as *const MyWaker)), // decrease refcount
    )
};

fn waker_into_waker(s: *const MyWaker) -> Waker {
    let raw_waker = RawWaker::new(s as *const (), &VTABLE);
    unsafe { Waker::from_raw(raw_waker) }
}

impl Task {
    fn new(reactor: Arc<Mutex<Reactor>>, data: u64, id: usize) -> Self {
        Task {
            id,
            reactor,
            data,
            is_registered: false,
        }
    }
}

impl Future for Task {
    type Output = usize;
    fn poll(mut self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output>
    {
        println!(
            "poll task {}, is_registered={},thread={:?}",
            self.id,
            self.is_registered,
            thread::current().id(),
        );
        let mut r = self.reactor.lock().unwrap();
        if r.is_ready(self.id) {
            Poll::Ready(self.id)
        } else if self.is_registered {
            r.insert_or_swap(self.id, cx.waker().clone());
            Poll::Pending
        } else {
            r.register(self.data, cx.waker().clone(), self.id);
            drop(r);
            self.is_registered = true;
            Poll::Pending
        }
    }
}

// ===== REACTOR =====
struct Reactor {
    dispatcher: Sender<Event>,
    handle: Option<JoinHandle<>>,
    readylist: Arc<Mutex<Vec<usize>>>,
    wakers: Arc<Mutex<HashMap<usize, Waker>>>,
}
#[derive(Debug)]
enum Event {

```

```

        Close,
        Timeout(u64, usize),
    }

impl Reactor {
    fn new() -> Self {
        let (tx, rx) = channel::<Event>();
        let readylist = Arc::new(Mutex::new(vec![]));
        let wakers: Arc<Mutex<HashMap<usize, Waker>>> =
Arc::new(Mutex::new(HashMap::new()));
        let rl_clone = readylist.clone();
        let wakers_clone = wakers.clone();
        let mut handles = vec![];
        let handle = thread::spawn(move || {
            // This simulates some I/O resource
            for event in rx {
                println!("REACTOR: {:?}" , event);
                let rl_clone = rl_clone.clone();
                let wakers_clone = wakers_clone.clone();
                match event {
                    Event::Close => break,
                    Event::Timeout(duration, id) => {
                        let event_handle = thread::spawn(move || {
                            thread::sleep(Duration::from_secs(duration));
                            rl_clone.lock().map(|mut rl| rl.push(id)).unwrap();
                            wakers_clone
                                .clone()
                                .lock()
                                .map(|mut wakers| wakers.remove(&id).map(|w|
w.wake()))
                                .unwrap();
                        });
                        handles.push(event_handle);
                    }
                }
            }

            for handle in handles {
                handle.join().unwrap();
            }
        });

        Reactor {
            readylist,
            dispatcher: tx,
            handle: Some(handle),
            wakers,
        }
    }

    fn insert_or_swap(&mut self, id: usize, waker: Waker) {
        self.wakers
            .lock()
            .map(|mut wakers| {
                let _ = wakers.insert(id, waker);
            })
    }
}

```



```
        })
        .unwrap();
    }

    fn register(&mut self, duration: u64, waker: Waker, data: usize) {
        self.insert_or_swap(data, waker);
        self.dispatcher
            .send(Event::Timeout(duration, data))
            .unwrap();
    }

    fn close(&mut self) {
        self.dispatcher.send(Event::Close).unwrap();
    }

    fn is_ready(&self, id_to_check: usize) -> bool {
        self.readylist
            .lock()
            .map(|rl| rl.iter().any(|id| *id == id_to_check))
            .unwrap()
    }
}

impl Drop for Reactor {
    fn drop(&mut self) {
        self.handle.take().map(|h| h.join().unwrap()).unwrap();
    }
}
```

Rust

## 完整的poll时wake问题

---

```

fn main() {
    let start = Instant::now();
    let reactor = Reactor::new();
    let future1 = Task::new(reactor.clone(), 1, 1);
    let future2 = Task::new(reactor.clone(), 2, 2);

    let fut1 = async {
        let val = future1.await;
        let dur = (Instant::now() - start).as_secs_f32();
        println!("Future got {} at time: {:.2}.", val, dur);
    };

    let fut2 = async {
        let val = future2.await;
        let dur = (Instant::now() - start).as_secs_f32();
        println!("Future got {} at time: {:.2}.", val, dur);
    };

    let mainfut = async {
        fut1.await;
        fut2.await;
    };

    block_on(mainfut);
    reactor.lock().map(|mut r| r.close()).unwrap();
}

use std::{
    collections::HashMap,
    future::Future,
    mem,
    pin::Pin,
    sync::{
        mpsc::{channel, Sender},
        Arc, Mutex,
    },
    task::{Context, Poll, RawWaker, RawWakerVTable, Waker},
    thread::{self, JoinHandle},
    time::{Duration, Instant},
};

// ===== EXECUTOR =====
fn block_on<F: Future>(mut future: F) -> F::Output {
    let mywaker = Arc::new(MyWaker {
        thread: thread::current(),
    });
    let waker = waker_into_waker(Arc::into_raw(mywaker));
    let mut cx = Context::from_waker(&waker);

    // SAFETY: we shadow `future` so it can't be accessed again.
    let mut future = unsafe { Pin::new_unchecked(&mut future) };
    let val = loop {
        match Future::poll(future.as_mut(), &mut cx) {
            Poll::Ready(val) => break val,
            Poll::Pending => {
                thread::sleep(std::time::Duration::from_millis(100));
            }
        }
    }
}

```

```

        println!("executor parked");
        thread::park()
    }
};

val
}

// ===== FUTURE IMPLEMENTATION =====
#[derive(Clone)]
struct MyWaker {
    thread: thread::Thread,
}

#[derive(Clone)]
pub struct Task {
    id: usize,
    reactor: Arc<Mutex<Box<Reactor>>>,
    data: u64,
}

fn mywaker_wake(s: &MyWaker) {
    let waker_ptr: *const MyWaker = s;
    let waker_arc = unsafe { Arc::from_raw(waker_ptr) };
    waker_arc.thread.unpark();
}

fn mywaker_clone(s: &MyWaker) -> RawWaker {
    let arc = unsafe { Arc::from_raw(s) };
    std::mem::forget(arc.clone()); // increase ref count
    RawWaker::new(Arc::into_raw(arc) as *const (), &VTABLE)
}

const VTABLE: RawWakerVTable = unsafe {
    RawWakerVTable::new(
        |s| mywaker_clone(&(s as *const MyWaker)), // clone
        |s| mywaker_wake(&(s as *const MyWaker)), // wake
        |s| mywaker_wake(*s as *const &MyWaker), // wake by ref
        |s| drop(Arc::from_raw(s as *const MyWaker)), // decrease refcount
    )
};

fn waker_into_waker(s: *const MyWaker) -> Waker {
    let raw_waker = RawWaker::new(s as *const (), &VTABLE);
    unsafe { Waker::from_raw(raw_waker) }
}

impl Task {
    fn new(reactor: Arc<Mutex<Box<Reactor>>>, data: u64, id: usize) -> Self {
        Task { id, reactor, data }
    }
}

impl Future for Task {
    type Output = usize;

```

```

fn poll(self: Pin<&mut Self>, cx: &mut Context<'_>) -> Poll<Self::Output> {
    let mut r = self.reactor.lock().unwrap();
    if r.is_ready(self.id) {
        println!("POLL: TASK {} IS READY", self.id);
        *r.tasks.get_mut(&self.id).unwrap() = TaskState::Finished;
        Poll::Ready(self.id)
    } else if r.tasks.contains_key(&self.id) {
        println!("POLL: REPLACED WAKER FOR TASK: {}", self.id);
        r.tasks
            .insert(self.id, TaskState::NotReady(cx.waker().clone()));
        Poll::Pending
    } else {
        println!(
            "POLL: REGISTERED TASK: {}, WAKER: {:?}" ,
            self.id,
            cx.waker()
        );
        r.register(self.data, cx.waker().clone(), self.id);
        Poll::Pending
    }
}

// ===== REACTOR =====
enum TaskState {
    Ready,
    NotReady(Waker),
    Finished,
}

struct Reactor {
    dispatcher: Sender<Event>,
    handle: Option<JoinHandle<()>>,
    tasks: HashMap<usize, TaskState>,
}

#[derive(Debug)]
enum Event {
    Close,
    Timeout(u64, usize),
}

impl Reactor {
    fn new() -> Arc<Mutex<Box<Self>>> {
        let (tx, rx) = channel::<Event>();
        let reactor = Arc::new(Mutex::new(Box::new(Reactor {
            dispatcher: tx,
            handle: None,
            tasks: HashMap::new(),
        })));

        let reactor_clone = Arc::downgrade(&reactor);
        let handle = thread::spawn(move || {
            let mut handles = vec![];
            for event in rx {
                let reactor = reactor_clone.clone();

```

```

        match event {
            Event::Close => break,
            Event::Timeout(duration, id) => {
                let event_handle = thread::spawn(move || {
                    thread::sleep(Duration::from_millis(duration));
                    let reactor = reactor.upgrade().unwrap();
                    reactor.lock().map(|mut r| r.wake(id)).unwrap();
                    println!("wake {}", id);
                });
                handles.push(event_handle);
            }
        }
    }
    handles
        .into_iter()
        .for_each(|handle| handle.join().unwrap());
});
reactor.lock().map(|mut r| r.handle = Some(handle)).unwrap();
reactor
}

fn wake(&mut self, id: usize) {
    self.tasks
        .get_mut(&id)
        .map(|state| match mem::replace(state, TaskState::Ready) {
            TaskState::NotReady(waker) => waker.wake(),
            TaskState::Finished => panic!("Called 'wake' twice on task: {}",
id),
            _ => unreachable!(),
        })
        .unwrap();
}

fn register(&mut self, duration: u64, waker: Waker, id: usize) {
    if self.tasks.insert(id, TaskState::NotReady(waker)).is_some() {
        panic!("Tried to insert a task with id: '{}', twice!", id);
    }
    self.dispatcher.send(Event::Timeout(duration, id)).unwrap();
}

fn close(&mut self) {
    self.dispatcher.send(Event::Close).unwrap();
}

fn is_ready(&self, id: usize) -> bool {
    self.tasks
        .get(&id)
        .map(|state| match state {
            TaskState::Ready => true,
            _ => false,
        })
        .unwrap_or(false)
}
}

```

```
impl Drop for Reactor {  
    fn drop(&mut self) {  
        self.handle.take().map(|h| h.join().unwrap()).unwrap();  
    }  
}
```

Rust

## 转载说明

---

本文允许转载,但是请注明出处.作者:stevenbai 本人博客:<https://stevenbai.top/>